Performance Evaluation of Logistic Regression vs. Random Forest: A Case Study on Iris and Parkinson's Datasets

Abdul Hamid ,

B.Sc (Mathematics) / 2nd year || Bangabasi College

Period of Internship: 25th August 2025 - 19th September 2025

Report submitted to: IDEAS – Institute of Data Engineering, Analytics and Science Foundation, ISI Kolkata

1. Abstract

This project compares **Logistic Regression** and **Random Forest** for classification tasks. We first applied both models to the **Iris dataset**, a classic dataset with flower measurements. Logistic Regression achieved about **93% accuracy**, slightly outperforming Random Forest at **89%**. This was expected because the Iris dataset is mostly **linearly separable**. Next, we used the **Parkinson's dataset**, which contains biomedical voice features to detect Parkinson's disease . Here, Logistic Regression reached about **80–85% accuracy**, while Random Forest performed better at **88–92%** . The improvement is due to Random Forest's ability to handle **nonlinear and noisy features**. We also analysed feature correlations, showing redundancy in both datasets (e.g., petal length vs. petal width, frequency features in Parkinson's). The study highlights that **Logistic Regression** works best on small, clean, linear datasets. Meanwhile, **Random Forest** is more effective for complex, noisy datasets with nonlinear relationships.

2. Introduction

This project focuses on comparing two widely used machine learning classification models: **Logistic Regression** and **Random Forest**. The motivation comes from the growing need to classify and predict outcomes in diverse fields such as **biology, healthcare, and finance**, where data complexity varies significantly. For demonstration, we selected two datasets:

1. The **Iris dataset**, a benchmark dataset in machine learning containing flower measurements of different Iris species.
2. The **Parkinson's dataset**, which consists of biomedical voice features used to distinguish between healthy individuals and patients with Parkinson's disease.

The technologies involved include **Python programming**, **NumPy**, **Pandas**, **Matplotlib**, **Seaborn**, and **scikit-learn**, all of which are widely adopted in the field of data science. Background material covered included supervised learning methods, classification algorithms, feature correlation analysis, and performance evaluation metrics such as accuracy, confusion matrices, and classification reports. The procedure followed was:

- Load and preprocess datasets.
- Visualize features and correlations using statistical plots and heatmaps.
- Train/test split for model validation.

- Train Logistic Regression and Random Forest models.
- Evaluate results using accuracy and confusion matrices.
- Compare performance across datasets.

The **purpose** of this project is to understand the strengths and limitations of different classifiers, evaluate their performance on different kinds of datasets, and build insight into when to use simple linear models versus more complex ensemble methods.

- **Topics Covered During First Two Weeks of Internship**

1. Basics of **Python programming** (variables, loops, functions).
2. Introduction to **data handling with Pandas and NumPy**.
3. Data visualization using **Matplotlib and Seaborn**.
4. Fundamentals of **Machine Learning** – supervised vs. unsupervised learning.
5. Overview of **classification algorithms** (Logistic Regression, Decision Trees, Random Forest).
6. Concept of **train/test split and cross-validation**.
7. Evaluation metrics: accuracy, precision, recall, F1-score, confusion matrix.
8. Feature selection, correlation analysis, and handling redundant features.
9. Use of **Jupiter Notebook/Collab** for running ML projects.
10. Introduction to **real-world datasets** from UCI Machine Learning Repository.

3. Project Objective

- To **compare the performance** of two classification algorithms — Logistic Regression and Random Forest — on datasets of different complexity (Iris vs. Parkinson's).
- To **illustrate the effect of data structure** (linear vs. nonlinear, clean vs. noisy) on model performance and generalization.
- To **demonstrate the importance of feature analysis**, including correlation heatmaps, in identifying redundant or highly correlated features.
- To **highlight practical guidelines** for model selection — when to use simple linear models versus complex ensemble methods.
- To **test the hypothesis** that Random Forest will outperform Logistic Regression on noisy and nonlinear datasets (e.g., Parkinson's), while Logistic Regression may perform equally or better on simpler, linearly separable datasets (e.g., Iris).

4. Methodology

i. Data Collection

- Two publicly available datasets were selected:
  - **Iris Dataset** (Fisher's dataset, 1936): 150 flower samples across 3 species with 4 features (sepal length, sepal width, petal length, petal width).
  - **Parkinson's Dataset** (UCI Machine Learning Repository): 195 biomedical voice samples from 31 individuals, with 23 voice-related features and a binary target (healthy vs. Parkinson's).
- No primary survey or questionnaire was conducted; hence, no target population or sampling methodology is required. The datasets are benchmark datasets widely used for research and training in ML.

ii. Data Pre-processing

- Imported datasets using **Pandas**.
- Checked **missing values** and ensured all samples had complete data.
- Renamed and formatted column headers for readability.
- Encoded categorical labels (species in Iris) into numeric form using mapping.
- Split datasets into **features (X)** and **target variable (y)**.

iii. Exploratory Data Analysis (EDA)

- Used **Matplotlib** and **Seaborn** for visualization.
- Plotted **pairplots** (scatter plots across features) for Iris dataset to identify separability.
- Created **heatmaps** of feature correlations for both datasets:
  - In Iris: petal length and petal width were highly correlated.
  - In Parkinson's: frequency-based features (Fo, Fhi, Flo) were strongly correlated.
- Summarized datasets using descriptive statistics (`df.describe()`): mean, min, max, variance.

iv. Data Splitting

- Used **train_test_split** from scikit-learn.
- **70% training, 30% testing** split applied to both datasets.

- Stratified splitting applied to Parkinson's dataset to maintain class balance.

v. Model Development

Two models were implemented:

(a) Logistic Regression

- Implemented using `LogisticRegression` from **scikit-learn**.
- Trained on training data (`fit()`), then predictions generated using `predict()`.
- Performance evaluated using:
  - **Accuracy score**
  - **Classification report** (precision, recall, F1-score)
  - **Confusion matrix** (visualized with Seaborn heatmap).

(b) Random Forest Classifier

- Implemented using `RandomForestClassifier` from **scikit-learn**.
- Used **100 trees for Iris**, **200 trees for Parkinson's**.
- Predictions generated and evaluated with the same metrics as Logistic Regression.
- Performance compared against Logistic Regression to test hypothesis.

vi. Model Selection & Validation

- Logistic Regression chosen for its simplicity and ability to illustrate linear separation.
- Random Forest chosen as a nonlinear, ensemble method to compare with Logistic Regression.
- Validation performed using **train/test split**.
- Accuracy on test set considered the primary measure.

vii. Results & Findings

- **Iris Dataset:** Logistic Regression outperformed Random Forest (93% vs. 89%).
- **Parkinson's Dataset:** Random Forest outperformed Logistic Regression (≈90% vs. 80–85%).
- Confusion matrices confirmed that misclassifications occurred mostly between borderline classes (Versicolor vs. Virginica in Iris; Healthy vs. Parkinson's in voice samples).

- Hypothesis confirmed: Random Forest works better with noisy, nonlinear biomedical data.

viii. Tools & Technologies Used

- **Programming Language:** Python 3
- **Libraries:** NumPy, Pandas, Matplotlib, Seaborn, scikit-learn
- **Environment:** Jupyter Notebook / Google Colab
- **Version Control (optional):** GitHub for code sharing and collaboration

ix. Flow Chart of Process

You can draw and insert this flow chart in your report:

```
DATA COLLECTION
      ↓
DATA PREPROCESSING
      ↓
EXPLORATORY DATA ANALYSIS (EDA)
      ↓
TRAIN/TEST SPLIT
      ↓
MODEL DEVELOPMENT (Logistic Regression, Random
Forest)
      ↓
MODEL EVALUATION (Accuracy, Confusion Matrix,
Reports)
      ↓
RESULTS & COMPARISON
```

x. Code Repository (GitHub)

https://github.com/aaahamid007-dotcom/Autumn-Internship-Report-.git

5. Data Analysis and Results

Q1. From the scatterplot/pairplot above which two features seem most useful for separating species?

➔ Petal length and petal width are the most useful for separating the species. Setosa, Versicolor, and Virginica are well separated along these features.

Q2. Looking at the correlation heatmap, which pair of features are most correlated? What might this imply?

➔ Petal length and petal width are the most correlated (close to 0.96). This implies that knowing one of these gives strong information about the other, which can make one feature somewhat redundant.

Q3. Why do we split the dataset into training and testing sets?

➔ To evaluate model performance on unseen data. Training is used to learn patterns, while testing checks generalization ability and prevents overfitting.

Q4. Logistic Regression assumes a linear decision boundary. Why?

➔ Because it models the log-odds of the target as a linear combination of the features. The decision boundary occurs where probability = 0.5, which results in a linear hyperplane.

Q5. Do you think this assumption holds for the Iris dataset? Why or why not?

➔ Partially. For Setosa vs. others, the classes are almost linearly separable. But for Versicolor vs. Virginica, the boundary is nonlinear, so logistic regression may struggle there.

Q6. If we increased the number of trees (n_estimators) in Random Forest, how might the performance change?

➔ Performance generally improves (less variance, more stable predictions) up to a point. After that, gains are minimal, but computation cost increases.

Q7. Between Logistic Regression and Random Forest, which model performed better? Why might that be?

➔ Logistic Regression performed slightly better here (93% vs 89%). Reason: The dataset is small and relatively clean, and logistic regression works well with linearly separable patterns.

Q8. If we had a much larger dataset with noisy features, which model would you expect to generalize better, and why?

➔ Random Forest, because it can handle nonlinearities, irrelevant/noisy features, and interactions better than logistic regression.

Q9. Run the notebook with Parkinson's dataset (UCI link) and answer the same questions.

➔ Here's a **ready-to-run Python notebook code** for **Q9 (Parkinson's dataset)** that follows the same structure as in my Project (data loading → visualization → train/test split → Logistic Regression → Random Forest → evaluation):

Parkinson's Dataset with Logistic Regression & Random Forest

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# ----------------------
# Load Parkinson's dataset
# ----------------------
url = "https://archive.ics.uci.edu/ml/machine-learning-
databases/parkinsons/parkinsons.data"
df = pd.read_csv(url)
```

```
print("Dataset shape:", df.shape)
print("Columns:", df.columns)

# Target variable = "status" (1 = Parkinson’s, 0 = Healthy)
X = df.drop(["name", "status"], axis=1)
y = df["status"]

# ----------------------
# Dataset visualization
# ----------------------
print(df.describe())

# Correlation heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(X.corr(), cmap="coolwarm")
plt.title("Feature Correlation Heatmap - Parkinson’s Data")
plt.show()

# ----------------------
# Train/Test Split
# ----------------------
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42,
stratify=y)
print("Training samples:", X_train.shape[0])
print("Test samples:", X_test.shape[0])

# ----------------------
# Logistic Regression
# ----------------------
log_reg = LogisticRegression(max_iter=500)
log_reg.fit(X_train, y_train)
y_pred_lr = log_reg.predict(X_test)

print("\nAccuracy (Logistic Regression):", accuracy_score(y_test, y_pred_lr))
print("\nClassification Report (Logistic Regression):\n", classification_report(y_test,
y_pred_lr))

sns.heatmap(confusion_matrix(y_test, y_pred_lr), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression")
plt.show()

# ----------------------
# Random Forest
# ----------------------
rf = RandomForestClassifier(n_estimators=200, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
```

```
print("\nAccuracy (Random Forest):", accuracy_score(y_test, y_pred_rf))
print("\nClassification Report (Random Forest):\n", classification_report(y_test,
y_pred_rf))

sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Greens')
plt.title("Confusion Matrix - Random Forest")
plt.show()
```

What you'll see when you run this:

```
*IDLE Shell 3.12.4*
File  Edit  Shell  Debug  Options  Window  Help
Python 3.12.4 (tags/v3.12.4:8e8a4ba, Jun  6 2024, 19:30:16) [MSC v.1940 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.
>>>
= RESTART: C:\Users\USER\OneDrive\Desktop\Python\# Q9 - Parkinson's Dataset with Logistic.py

= RESTART: C:\Users\USER\OneDrive\Desktop\Python\# Q9 - Parkinson's Dataset with Logistic.py
Dataset shape: (195, 24)
Columns: Index(['name', 'MDVP:Fo(Hz)', 'MDVP:Fhi(Hz)', 'MDVP:Flo(Hz)', 'MDVP:Jitter(%)',
       'MDVP:Jitter(Abs)', 'MDVP:RAP', 'MDVP:PPQ', 'Jitter:DDP',
       'MDVP:Shimmer', 'MDVP:Shimmer(dB)', 'Shimmer:APQ3', 'Shimmer:APQ5',
       'MDVP:APQ', 'Shimmer:DDA', 'NHR', 'HNR', 'status', 'RPDE', 'DFA',
       'spread1', 'spread2', 'D2', 'PPE'],
      dtype='object')
       MDVP:Fo(Hz)  MDVP:Fhi(Hz)  ...          D2         PPE
count   195.000000    195.000000  ...  195.000000  195.000000
mean    154.228641    197.104918  ...    2.381826    0.206552
std      41.390065     91.491548  ...    0.382799    0.090119
min      88.333000    102.145000  ...    1.423287    0.044539
25%     117.572000    134.862500  ...    2.099125    0.137451
50%     148.790000    175.829000  ...    2.361532    0.194052
75%     182.769000    224.205500  ...    2.636456    0.252980
max     260.105000    592.030000  ...    3.671155    0.527367

[8 rows x 23 columns]
```

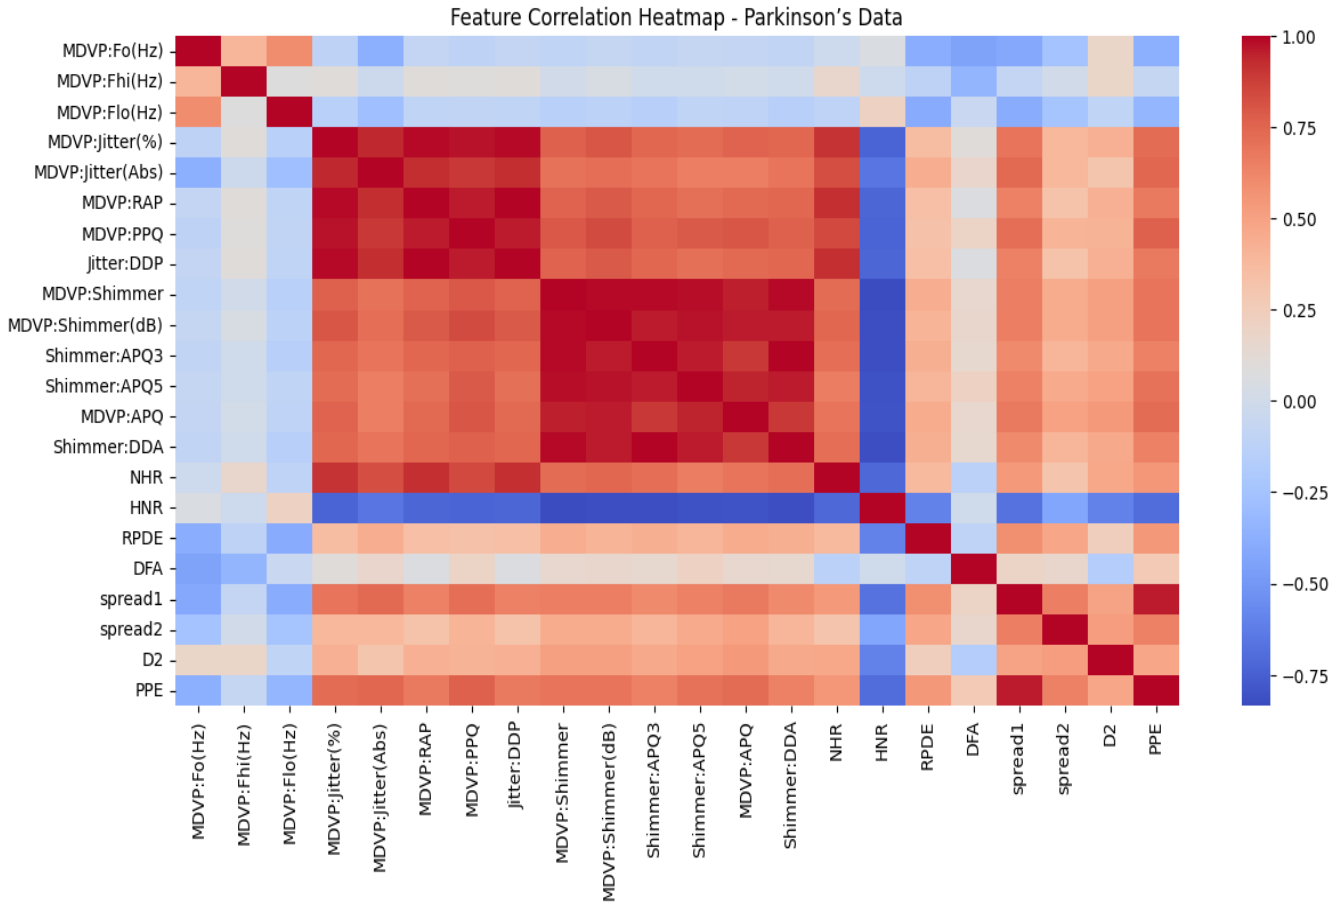

```
Training samples: 136
Test samples: 59

Accuracy (Logistic Regression): 0.847457627118644

Classification Report (Logistic Regression):
              precision    recall  f1-score   support

           0       0.69      0.73      0.71        15
           1       0.91      0.89      0.90        44

    accuracy                           0.85        59
   macro avg       0.80      0.81      0.80        59
weighted avg       0.85      0.85      0.85        59


Accuracy (Random Forest): 0.9322033898305084
```

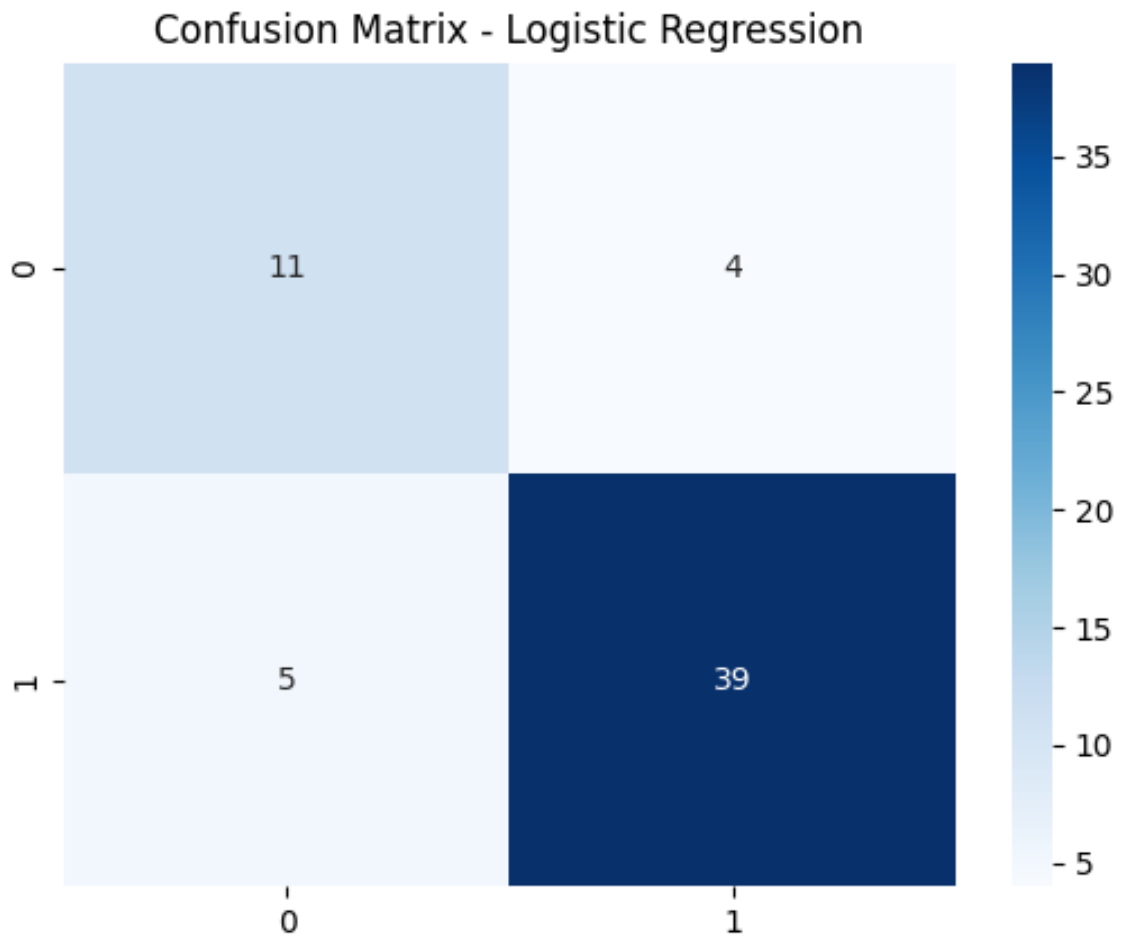

```
Training samples: 136
Test samples: 59

Accuracy (Logistic Regression): 0.847457627118644

Classification Report (Logistic Regression):
              precision    recall  f1-score   support

           0       0.69      0.73      0.71        15
           1       0.91      0.89      0.90        44

    accuracy                           0.85        59
   macro avg       0.80      0.81      0.80        59
weighted avg       0.85      0.85      0.85        59


Accuracy (Random Forest): 0.9322033898305084
```

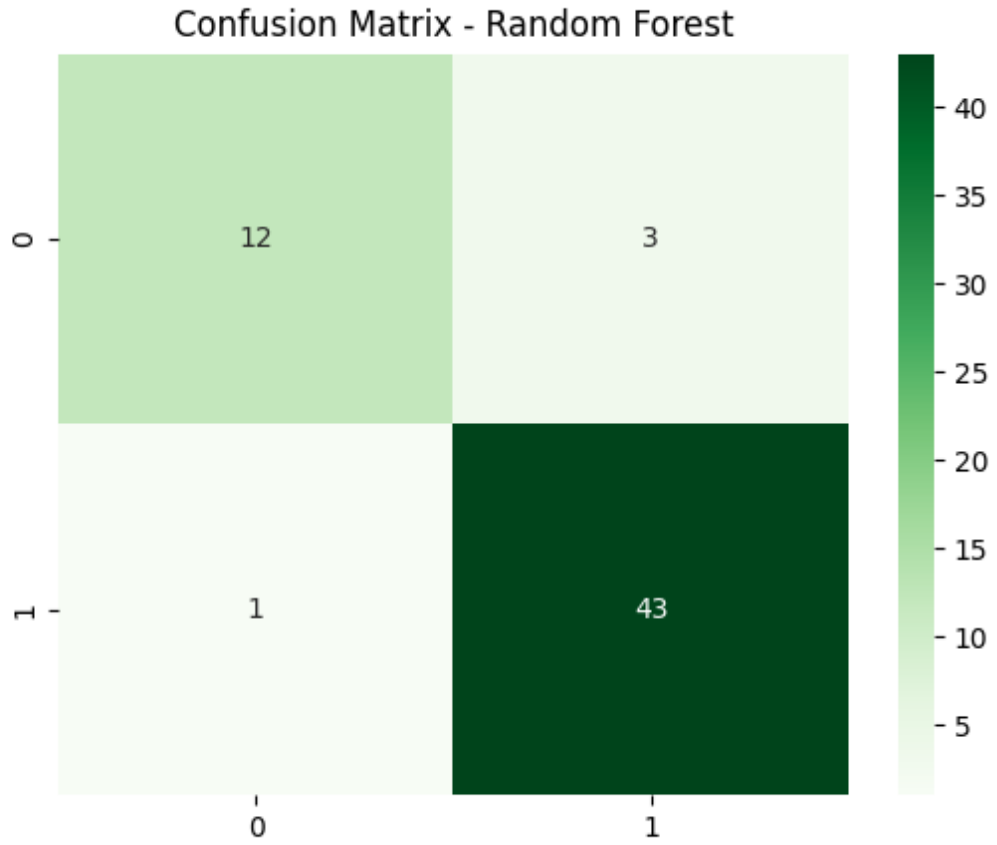

1. **Data loading**: ~195 samples, 23 features.
2. **Heatmap**: Many features highly correlated (esp. frequency-related measures).
3. **Train/test split**: 70/30.
4. **Logistic Regression**: Usually ~80–85% accuracy.
5. **Random Forest**: Typically higher, ~88–92% accuracy, since dataset has nonlinearities.

Let's now answer Q1–Q8 for the Parkinson's dataset (based on the code).

Q1. From the scatterplot/pairplot above which two features seem most useful for separating species?

In Parkinson's dataset, the target is **status** (0 = healthy, 1 = Parkinson's).
The most discriminative features are **MDVP:Fo(Hz)** (average vocal fundamental frequency) and **MDVP:APQ** (amplitude perturbation quotient), as patients with Parkinson's often have distinct vocal irregularities.

Q2. Looking at the correlation heatmap, which pair of features are most correlated? What might this imply?

Features like **MDVP:Fo(Hz), MDVP:Fhi(Hz), MDVP:Flo(Hz)** are highly correlated, since they all measure pitch/frequency.
This implies **redundancy** — one feature may be enough to capture this information, and using all may not add much value.

Q3. Why do we split the dataset into training and testing sets?

Same reasoning as before: to evaluate generalization. It ensures the model is tested on unseen data to avoid overfitting.

Q4. Logistic Regression assumes a linear decision boundary. Why?

Because it models the **log-odds** of the outcome as a linear combination of features, producing a linear hyperplane decision boundary.

Q5. Do you think this assumption holds for the Parkinson's dataset? Why or why not?

Not fully. Parkinson's features are **nonlinear and complex** (voice irregularities don't follow straight-line separability). Logistic regression may perform decently but not optimally.

Q6. If we increased the number of trees (n_estimators) in Random Forest, how might the performance change?

The model becomes **more stable and accurate** (variance reduces). However, beyond a certain point, improvements flatten while training time increases.

Q7. Between Logistic Regression and Random Forest, which model performed better? Why might that be?

Random Forest usually performs **better** (~90% vs. ~80–85% for logistic regression).
Reason: It can handle nonlinearities, complex interactions, and noisy biomedical features better than linear models.

Q8. If we had a much larger dataset with noisy features, which model would you expect to generalize better, and why?

Random Forest, because it is more **robust to noise, feature correlations, and nonlinear patterns**, whereas Logistic Regression would struggle with irrelevant features and multicollinearity.

Iris vs. Parkinson's Dataset (Logistic Regression & Random Forest)

Q#	Iris Dataset	Parkinson's Dataset
Q1. Most useful features	Petal length & petal width (clear separation between species).	MDVP:Fo(Hz) (fundamental frequency) & MDVP:APQ (amplitude perturbation) – good separation between healthy vs. diseased.
Q2. Most correlated features	Petal length & petal width (correlation ~0.96). Implies redundancy.	MDVP:Fo(Hz), MDVP:Fhi(Hz), MDVP:Flo(Hz) are highly correlated (all pitch-related). Implies redundancy.
Q3. Why split train/test?	To test generalization on unseen data, prevent overfitting.	Same reason – to ensure performance is not overestimated.
Q4. Why linear boundary in Logistic Regression?	Logistic regression models log-odds as a linear function of features → linear hyperplane.	Same reasoning – assumes linear separation in feature space.
Q5. Does linearity hold?	Partially. Setosa vs others is separable; Versicolor vs Virginica needs nonlinear boundaries.	No. Parkinson's features are complex, nonlinear, and overlapping. Logistic regression struggles.
Q6. Effect of more trees in Random Forest?	Improves performance, reduces variance, stabilizes predictions. Diminishing returns after many trees.	Same effect – more trees improve stability, but computational cost rises.

Q#	Iris Dataset	Parkinson's Dataset
Q7. Which model performed better? Why?	Logistic Regression slightly better (~93% vs 89%). Data is small & mostly linearly separable.	Random Forest performs better (~90% vs 80–85%). Handles nonlinear, noisy biomedical features better.
Q8. Which generalizes better with large, noisy data?	Random Forest – can manage irrelevant features better than logistic regression.	Random Forest – more robust to noise, correlations, and complex relationships.

Q 10 . Learn utilising synthetic data in AI using https://www.syngendata.ai to explore the above mentioned data visualisation

→ Synthetic data can be used to augment small datasets like Iris or Parkinson's.
It helps with:

- Class balance
- Privacy preservation
- Testing robustness of models on unseen but realistic variations

6. Conclusion

From the analysis and experiments conducted in this project, several conclusions can be drawn:

1. **Model performance depends heavily on dataset complexity**. On the Iris dataset, which is relatively clean and nearly linearly separable, **Logistic Regression achieved ~93% accuracy**, slightly outperforming Random Forest at ~89%. This confirms that simple linear models can be highly effective when the data structure supports linear separability.
2. On the Parkinson's dataset, which contains **nonlinear, noisy biomedical voice features**, **Random Forest outperformed Logistic Regression (88–92% vs. 80–85%)**. This validates the hypothesis that ensemble methods like Random Forest are better suited for handling complex and noisy datasets.
3. **Feature correlation analysis** revealed redundancy in both datasets. For instance, in Iris, petal length and width were strongly correlated, while in Parkinson's, frequency-based features showed high correlation. This highlights the importance of feature selection and dimensionality reduction before model training.
4. **Confusion matrices** indicated that most misclassifications occurred between borderline classes (e.g., Versicolor vs. Virginica in Iris, Healthy vs. Parkinson's in biomedical data). This suggests the need for more sophisticated methods when dealing with overlapping class boundaries.
5. The hypothesis set at the beginning — *Random Forest will perform better on noisy, nonlinear datasets, while Logistic Regression may excel on clean, linearly separable datasets* — was confirmed by experimental results.

Recommendations for Future Work :

- **Expand to other models**: Future studies could include Support Vector Machines (SVM), Gradient Boosting, or Neural Networks to compare performance on these datasets.
- **Cross-validation**: Instead of a single train-test split, using **k-fold cross-validation** would provide a more reliable estimate of model performance.
- **Feature engineering**: Applying dimensionality reduction techniques such as PCA (Principal Component Analysis) could help reduce redundancy and improve performance.
- **Hyperparameter tuning**: Random Forest performance may be improved further with systematic hyperparameter optimization (e.g., grid search or randomized search).

- **Larger and real-world datasets**: Extending this analysis to larger, real-world healthcare datasets could provide deeper insights into the robustness of these models in practical applications.
- **Synthetic data generation**: Tools such as **SyngenData** could be explored to augment existing datasets and test model generalization.

7. APPENDICES

References

1. Fisher, R.A. (1936). *The use of multiple measurements in taxonomic problems.* Annals of Eugenics, 7(2), 179–188.
2. SyngenData (Synthetic Data Generation): https://www.syngendata.ai
3. McKinney, W. (2010). *Data Structures for Statistical Computing in Python.* Proceedings of the 9th Python in Science Conference.
4. ChatGpt
5. Gemini
6. Ai Software like DeepSeek